

Predicting Blue-Green Algae Quantity in a Semi-Urban Reservoir

Nathan Jahns, M.Sc.

Data Science Intensive Capstone Project, Sept 23 Cohort




Background: Blue-Green Algae

- Harmful algae blooms (HABs), mostly blue-green algae, in lakes across USA every summer and can grow over a few days
- Can release toxins which kill fish, waterfowl, livestock, and pets and make humans very ill
- Measure phycocyanin (pigment in blue-green algae) with *in situ* probe
- Lake operators monitor readings and shut down recreation
 - Immediate, no warning to the public or staff, management actions are reactionary
 - Method to predict blue-green algae bloom before a shutdown would allow for planning and allocation of resources



Objective

- Develop machine learning models to predict phycocyanin using water quality measurements at a semi-urban reservoir I named Lake Fictitious
- Establish models capable of forecasting phycocyanin multiple days in advance using the explanatory variables collected at the reservoir
- With this model and data the magnitude of a HAB could be predicted before it occurs which would avoid immediate shutdowns



Phycocyanin

Data Acquisition

- Data collected at Lake Fictitious every 30 minutes for three years
- *In situ* AlgaeTracker manufactured by AquaRealTime
- 50,071 rows and 20 columns




Data Cleaning

- Removed unneeded columns
 - Phycocyanin, sunlight, water temperature, water turbidity, and chlorophyll-a
- Removed duplicate date and time
 - Calculated mean because column values for pairs were often different
 - Both forward and backward directions from each NaN value
- Inserted missing half hour rows
 - Largest gaps was 6.25 days, not too large to impute

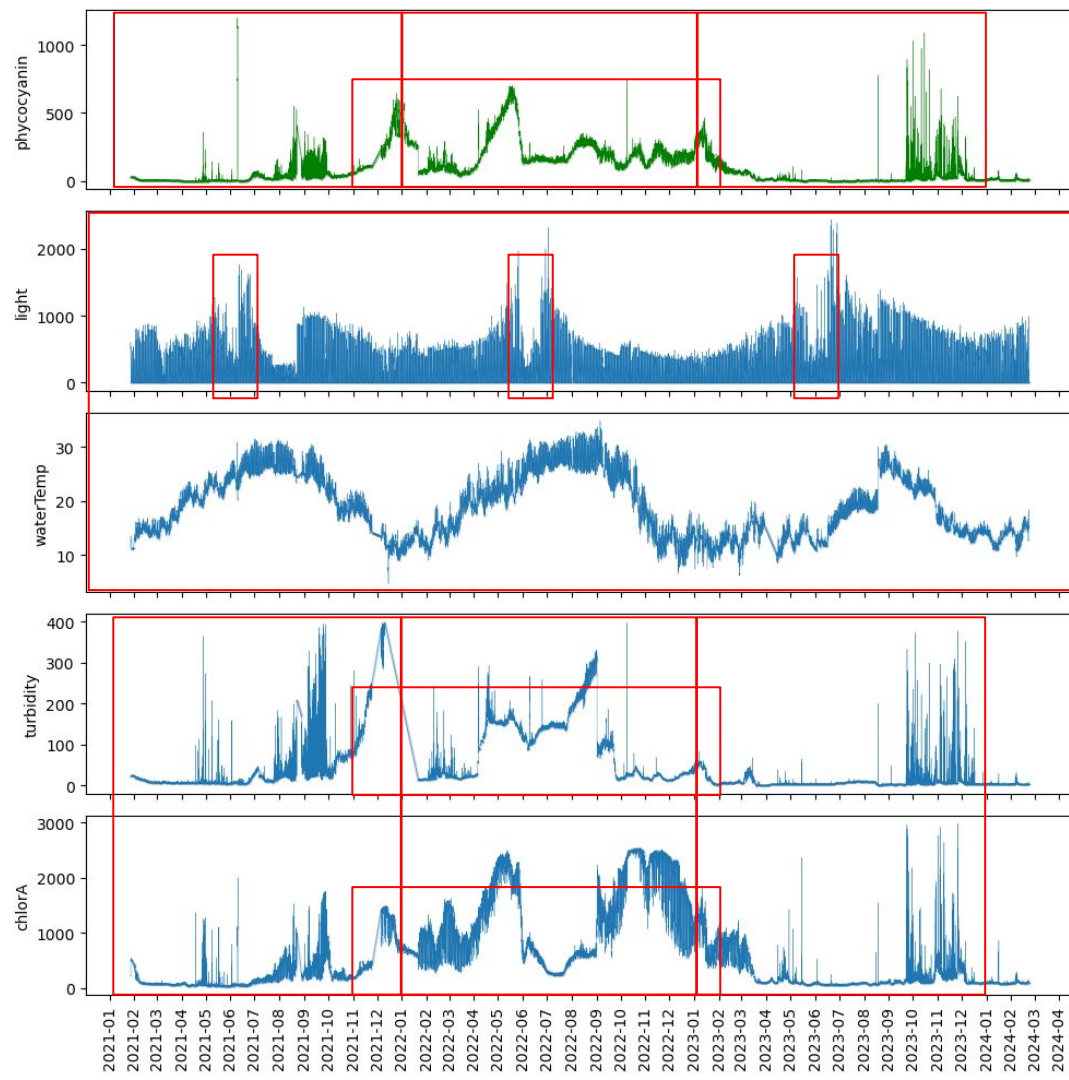
Data Cleaning

- Unreliable data
 - Readings were collected every hour
 - Low sunlight readings
 - Unit flipped upside
 - Water temperature unlikely to be 0 °C.
 - Maximums also unlikely.
- Replaced with NaNs
- Sunlight - replaced NaNs with mean 2022 and 2023 value
- Spline linear interpolation to impute elsewhere



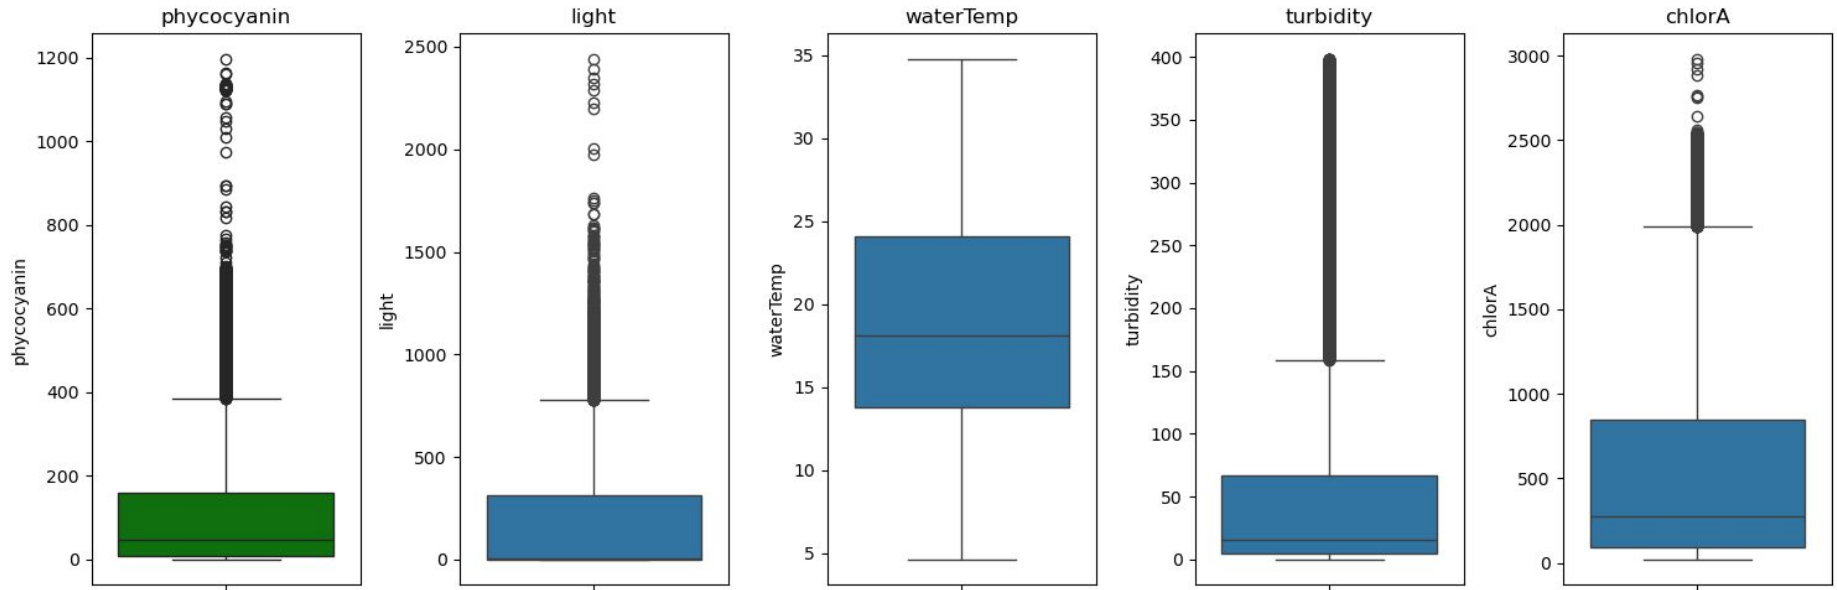
EDA: Trends & Patterns

- Sunlight and water temperature were seasonal
- June was cloudy each year and 2021 was also cloudy in July, Aug, and Dec
- Phycocyanin, turbidity, and chlorophyll-a not seasonal and three years were different
- Baseline was higher and readings often did not return to baseline for all metrics from 11/21-2/23








EDA: Distribution

- Water temperature fell within the whiskers (1.5 times the IQR)
- Other values greater than 3 std from the median were not removed because no reason to believe these are erroneous.



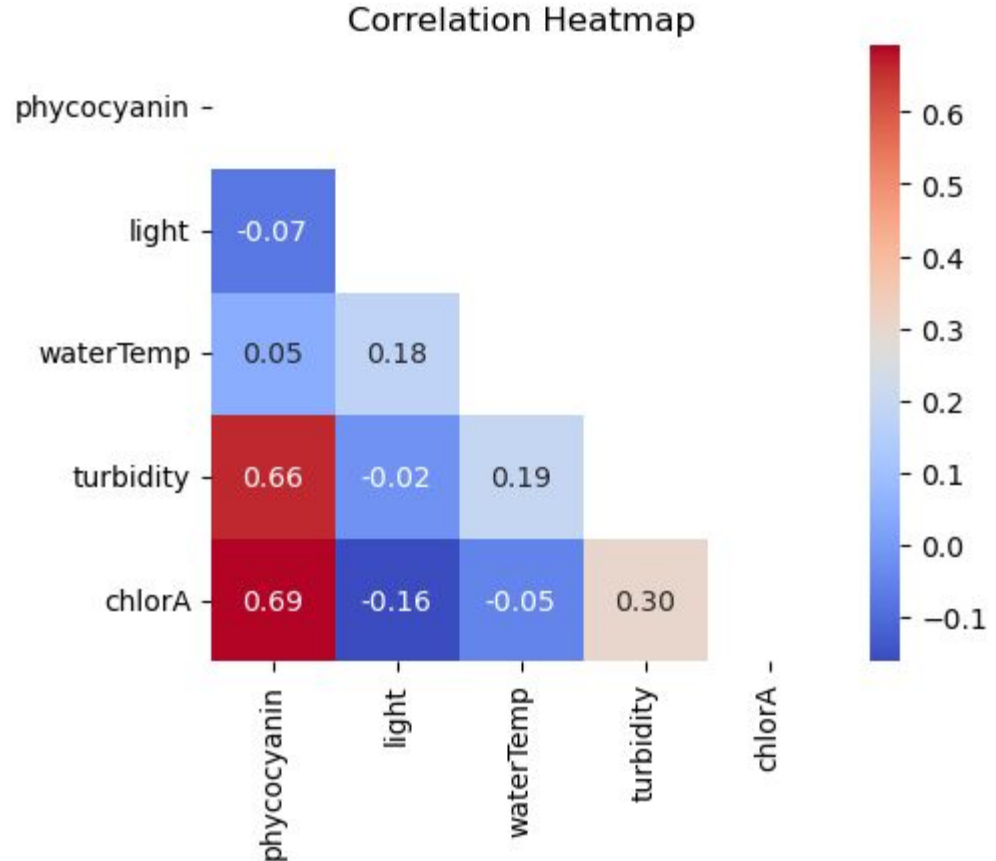
EDA: Distribution

- Reflec boxplots
- Not normally distributed
- Right-skewed




EDA: Correlations

- Moderately strong correlations between phycocyanin and turbidity and between phycocyanin and chlorophyll-a.
- Moderately weak correlations between turbidity and chlorophyll-a.
- Metrics correlated with phycocyanin likely to be good predictors
- Metrics with a low correlation to phycocyanin likely to be less relevant in the model



EDA: Stationarity

- KPSS and ADF tests for each metric
- Temperature, turbidity, and chlorophyll-a metrics are non-stationary
- Phycocyanin and sunlight
 - KPSS test indicated non-stationary
 - ADF test indicated stationary
 - “Difference stationary”
- All metrics need to be detrended



Engineering

- One day mean of metrics so same as predictions
 - 1,123 rows
- Reran Shapiro-Wilks tests
 - Not normal distributions
- Calculated natural logarithms
- Reran Shapiro-Wilks tests
 - Normal distributions
- Reran KPSS and ADF tests
 - Results similar to before
- Detrended with differencing to make stationary
- Reran KPSS and ADF tests
 - All metrics stationary
- Created 1 through 20 day lagged versions of the original features and the logged and differenced features




Pre-processing

- Defined two different sets of X and y
 - Original data and lagged features of the original data
 - Transformed features and lagged features of the transformed data
- Allowed comparison of the performance of the models using unaltered and transformed data

Modeling

- Regression, random forest, XGBoost
- Cross-validation using TimeSeriesSplit
- Root mean squared error (RMSE)
- Models were run multiple times starting with including non lagged feature data and adding a level of lagged feature data with each iteration up through the 20 lag features
- Performed for both sets of X and y
- RMSEs for each iteration were graphed and model with lowest RMSE was identified
- Lowest RMSE from models using transformed data were un-transformed before comparison.



Results

- 0.7-3.3% of the phycocyanin range (1654.64) which is considered low error.
- Median (55) error would not constitute a blue green bloom at a dangerous level
- Hypertuning for best RMSE of 12.4



	Original Data		Transformed Data	
Model	Lagged Features	RMSE	Lagged Features	RMSE
Linear Regression	0-3	29.5	0	14.4
Random Forest	0-13	43.7	0-8	12.8
XGBoost	0-14	55.3	0-10	12.5

Conclusion

- XGBoost model using transformed data lagged at zero to 11 days produced the highest performing model to predict phycocyanin in Lake Fictitious
- Error is small in the grand scheme of HABs where phycocyanin values are, at a minimum, in the several hundreds
- Model could be very helpful for predicting blue-green algae in a lake before implementing a shutdown, allowing for planning and allocation of resources



Special Thank You!

To Chris Lee at AquaRealTime for sharing data and answering many questions!

